

# LANCER Site Visit



# Outline

- [15 minutes] Introduction (Nate & Wen)
  - Team Introductions
  - Technical Approach
- [10 minutes] Progress Since Kick-Off (Nate & Wen)
  - Executive Summary
  - Planned Trajectory for end of Phase I
- [15 minutes] Collaboration Efforts (Nate & Rebecca)
  - CAGE
  - Talking to Kryptowire
  - Network Action Space
- [60 minutes] Early Results
  - [20 minutes] NetKAT (Jules & Nate)
  - [30 minutes] Inverse RL (Nico/Rebecca & Wen)
  - [10 minutes] Aether: Pronto + OnRamp (Hussain & Nate)
- [20 minutes] Response to Crawl Questions (Everyone)
- [30 minutes] Budget & Contracting (Shailja & Nate)

# Introduction



# Progress



# Progress

- Got going with Kryptowire TA1 Platform
- Started development using CAGE 2
- Started Modeling Red Agents Using Inverse RL
- Fast NetKAT implementation
- Standing Up Aether OnRamp

# Trajectory

- Crawl (6 month)
- Walk (6 month)
- Run (6 month)

# Collaboration Efforts



# Cage Challenge Overview

- Scenario of a network attack
- Goal: Blue Agent (defensive) stops the Red Agent (malicious) without disrupting the Green Agent (normal users)
- Integrated with CybORG, a reinforcement learning gym




Figure 1: Network of the scenario and challenge problem (Cage Challenge 2)



# Red Agent Actions

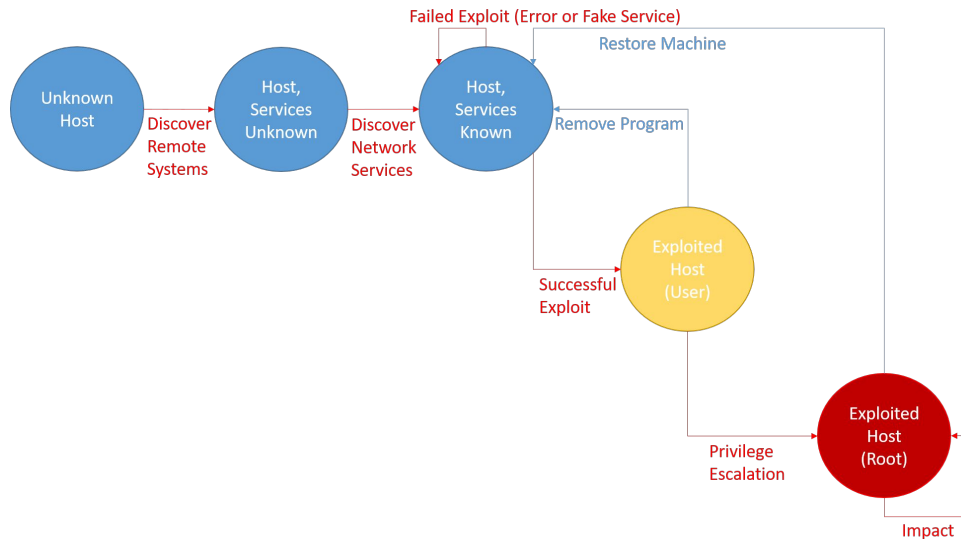
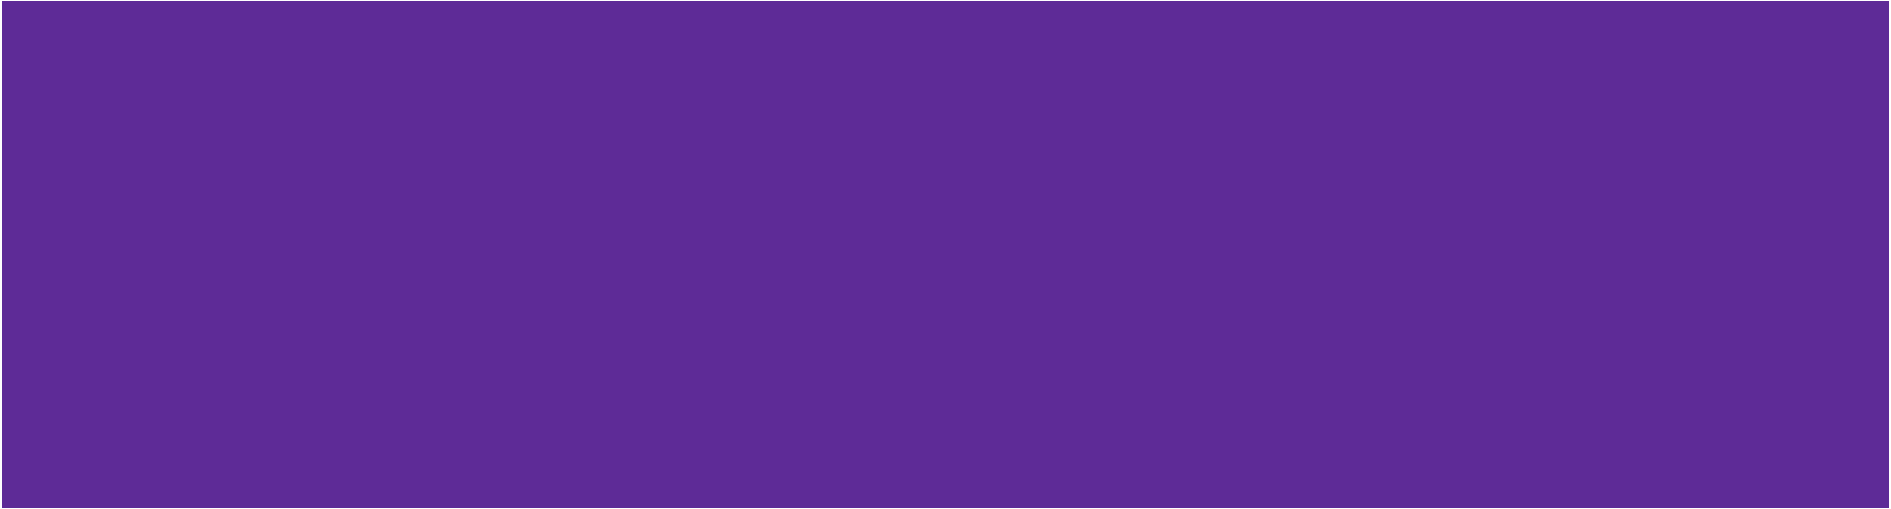


Figure 2: Effect of actions on host state (Cage Challenge 2)

# Early Results



# RL Outline

- Platform - describe cage challenge
- Dataset
- Model baselines (BC/SL has these limitations...)
- IRL results
  - Focus on IRL, no progress on RL side yet
- Plan moving forward
- Tables with data, add graphical aids

# Discuss overall approach

Why are we using imitation learning to learn the red agent

# Definitions

Define RL terminology: states, actions, observations

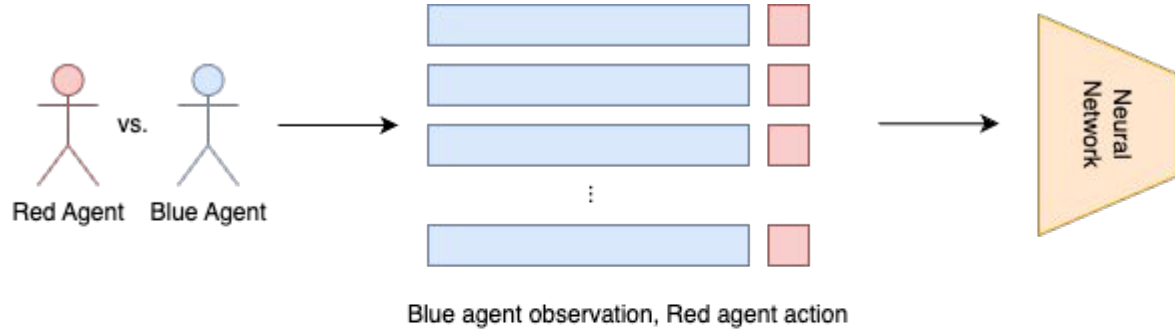
# Imitation Learning Approaches

1. Behavior Cloning (BC)
  - a.
2. Generative Adversarial Imitation Learning (GAIL)
  - a.

# Behavior Cloning (BC)

1. Collected data from environment with Blue, Green, Red agents
  - a. (Blue agent observation, Red agent action)
2. Used data to train a neural network to predict the Red agent action from the Blue agent observation
3. Created a Red agent that used this neural network to determine the next action (learned agent)
4. Collected reward from environment with a Blue agent, Green agent, and the learned agent

# Behavior Cloning






# Generative Adversarial Imitation Learning (GAIL)

**Discuss distribution shift, use self driving visual examples**

# Results



# Detailed results for IRL

**Plot showing BC improvement with increasing  
number\_of\_states input**

# BC: 0 Prev. States, 50 eval. Episodes

Red Agent	Blue Agent	Steps per Episo de	Dataset Collecti on Episode s	Train Loss	Train Accuracy	Validation Accuracy	Average Reward (50 trials)	Standard Deviation	Real Agent Average Reward (3 runs)	Real Agent Average Std. Dev.
B-Line	React Remove	100	100	0.16	0.95	0.93	556	361	947	193
B-Line	React Restore	100	100	0.64	0.77	0.77	-10.0	0.0	508	366
Meander	React Remove	100	100	0.71	0.72	0.67	11.1	39.5	630	259
Meander	React Restore	100	100	1.1	0.56	0.53	3.55	7.77	185	210

# BC: 3 Prev. States, 50 eval. Episodes

Red Agent	Blue Agent	Steps per Episode	Dataset Collection Episodes	Train Loss	Train Accuracy	Val Accuracy	Average Reward (50 trials)	Standard Deviation	Real Agent Average Reward (3 runs)	Real Agent Average Std. Dev.
B_Line	React Remove	100	100	0.038	0.986	0.967	693.575333 3	304.785803	946.573333 3	192.988038
B-Line	React Restore	100	100	0.0372	0.987	0.965	483.727333 3	335.984490 7	508.29	366.422970 1
Meander	React Remove	100	100	0.327	0.870	0.710	254.538666 7	245.967610 3	630.17	258.94
Meander	React Restore	100	100	0.615	0.762	0.587	77.038666 7	141.286532	185.01	209.86

# Next Steps

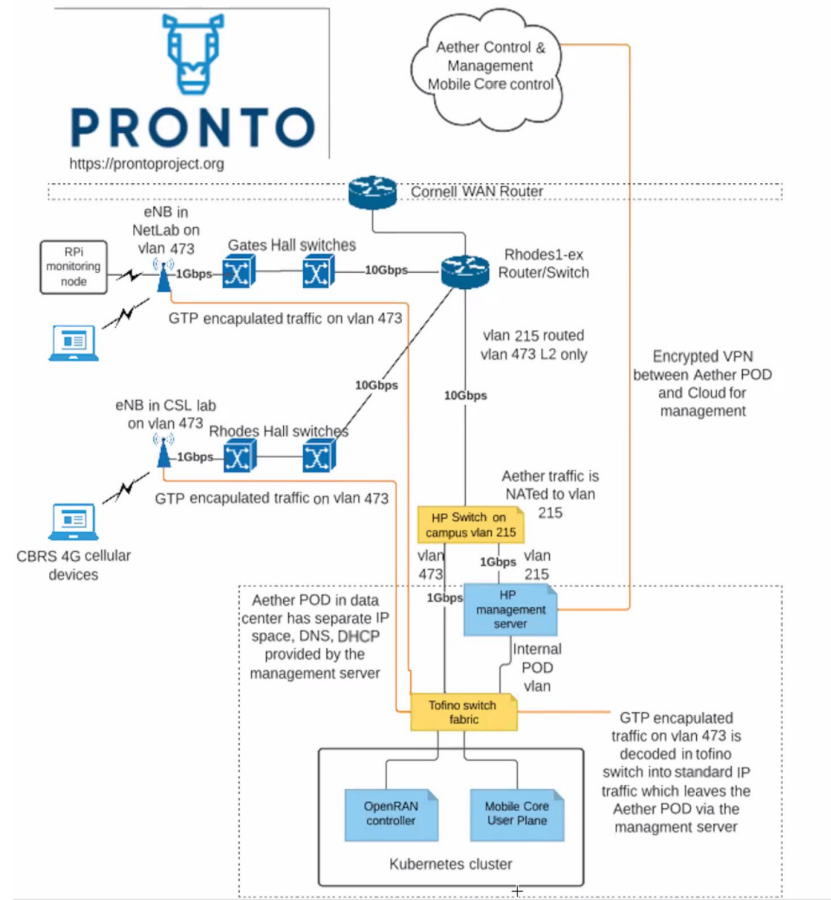
1. New IRL algorithms for improving modeling red agents;
2. Training RL agents against the learned red agents



# NetKAT

# Pronto Cornell Network

## Previous Cornell Network



# AetherOnRamp

Private Enterprise 5G network

- operational cluster that is capable of running 24/7 and supports live workloads on Cornell Network.

Mobile Network two main subsystems :

1. RAN - manages radio resources(spectrum)
2. Mobile Core - provide packet data network to mobile subscribers

# Aether

## Pronto pods Demo

- Demo setup with the end to end connectivity
  - End to end connectivity, Raspberry PI reachable from the Intel server
  - Grafana Dashboards
    - Gov for traffic, syncing from the UEs
- Working Demo -
  - MotoG phone connecting to the Aether APN(Access Point Name) within Pronto Network

## Aether OnRamp Progress

- Demo the setup, ideally with UEs and end to end connectivity
  - Current issues - Radio connectivity not syncing with the local Aether Core
  - Demo Setup - Grafana Dashboards

# Establish Secure Channel and Communication


Each SIM has global identifier :

- IMSI - unique 15 digit code

UE registration with Core :

- Communicates with nearby station
- Base station forwards request to the control plane, only gets auth relevant mapping in Core.

UE traffic routed using user plane



# Crawl Questions

- Learn about one another's approaches, find integration points, and collaborate on shared infrastructure
- What network should we model first and what workflows should be present?
- What agent actions will be simulated and executed?
- What is a 'good' resiliency criteria and how will we judge whether your approach is successful?
- What data types are needed for each performer and what data can be provided by each performer?
  - Data for attackers
  - Reward function for defenders (domain knowledge, Inverse RL)
- How do we collaborate on API design and code interfaces?
- What open-source technology can enable an end-to-end integration demo quickly?
- Who is the intended operator of your approach and what is the desired impact/benefit to their job?